

Inverse Simulation Tomography

Jannis.Maron@uni-siegen.de

February 13, 2026

Foreword

This document should serve as a guide on how to run each task in the *Inverse Simulation Tomography* project.

Within this project, we implemented three tasks:

1. **Forward Simulation:** Simulate multi-scattering wave propagation based on a given 3D Refractive Index (RI) distribution and a sequence of predefined poses.
2. **Pose Optimization:** Find poses for our *Forward Simulation* that best match the phase and amplitude to a recorded video sequence.
3. **Refractive Index Optimization:** Refine the estimated RI distribution by comparing the optimized poses to the ground truth video sequences.

A description of how to run each task and what outputs we expect can be found in the corresponding section.

Note that this guide does not aim to explain any technical details.

A high-level schematic overview for the combination of Pose and RI optimization can be found in Fig. 1.

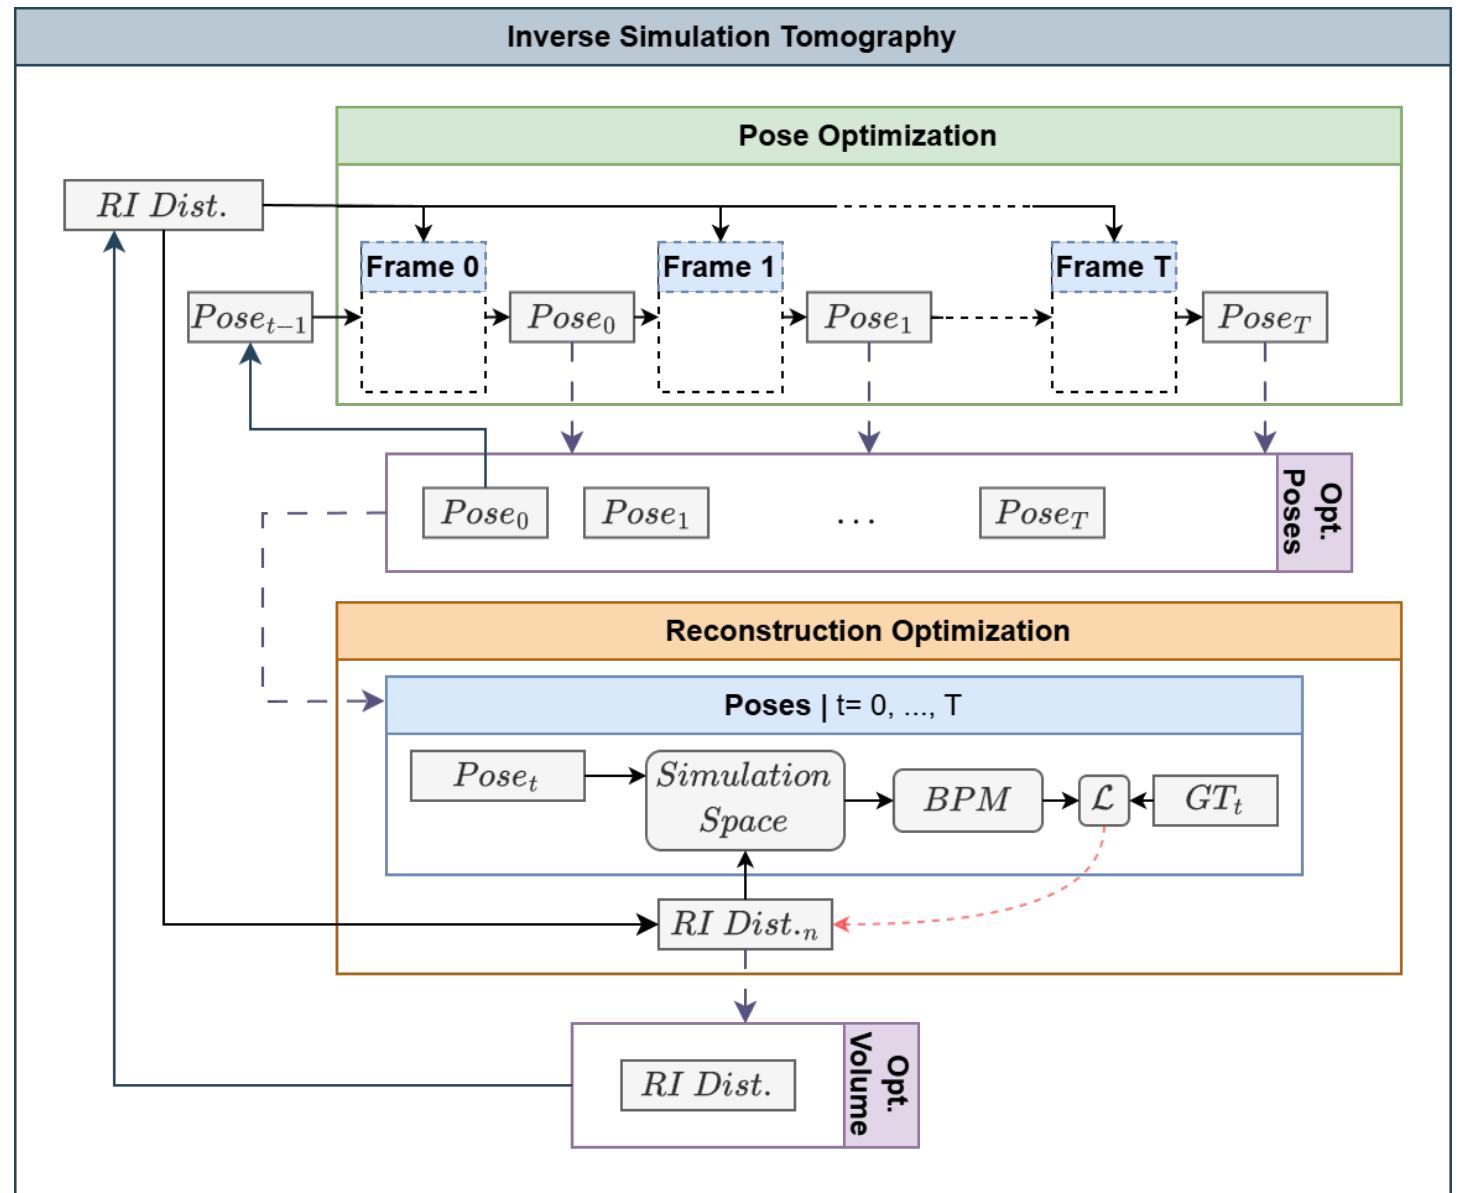


Figure 1: Top-level Overview

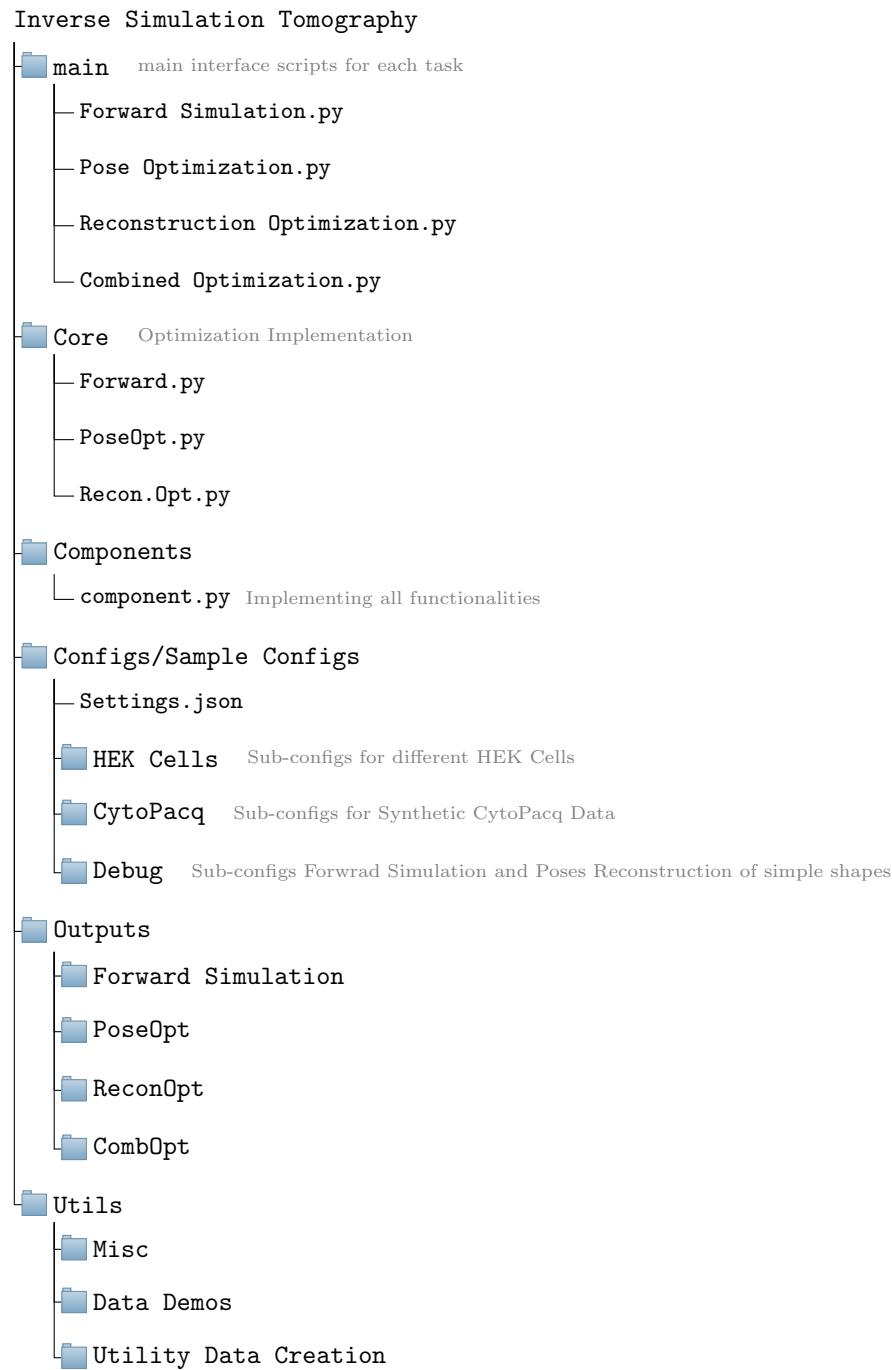


Figure 2: Assumed Project Hierarchy

Contents

1 Settings	5
2 Sub-configs	6
2.1 Simulation Config.json	6
2.2 Data Config.json	7
3 Forward Simulation	8
3.1 Configs	8
3.1.1 General Settings	8
3.1.2 Forward Config.json	9
3.2 Outputs	10
4 Pose Optimization	12
4.1 Configs	13
4.1.1 General Settings	13
4.1.2 PoseOpt Config.json	13
4.2 Outputs	15
5 Reconstruction Optimization	18
5.1 Configs	18
5.1.1 General Settings	18
5.1.2 ReconOpt Config.json	19
5.2 Outputs	20
6 Combined Optimization	22
6.1 Configs	23
6.1.1 General Settings	23
6.2 Outputs	23
7 Domain Adaptation / Post-Processing Transforms	24
8 Custom Data	26
8.1 Ground Truth Data	26
8.2 Voxel Object	26
8.3 Gradient Masks	27

1 Settings

Each task is configured through a dedicated settings file. If any main script is executed without specifying a settings file, the user will be prompted to select one interactively. The setting file is structured .json documents and consists of two main blocks:

- a **Common** block containing settings shared across all tasks,
- a **Task** block containing task-specific parameters.

Task blocks are described in their corresponding sections¹.

Common Settings

Example: Common Settings Block

```
{
  "Common": {
    "config_dir": "../Configs/Sample_Configs",
    "device": null,
    "dtype": "float32",
    "phase_unwrap": false
  }
}
```

Parameter Description

config_dir

Path relative to the executing script that points to the root directory containing all sub-configuration files.

device

Compute device selection.

- `null`: Automatically select GPU if available, otherwise CPU
- `"gpu", "cuda"`: Force GPU execution
- `"cpu"`: Force CPU execution

dtype

Floating-point precision used throughout the computation.

- `float32`
- `float64`

Warning: Higher precision significantly increases memory usage.

phase_unwrap

Boolean flag controlling wrapped versus unwrapped phase visualization. This setting affects visualization only and has no impact on computation.

¹Reused sub-settings are documented in separate sections.

2 Sub-configs

For improved readability and reusability, configuration parameters that belong to distinct sub-tasks are stored in separate *sub-config* files. Each sub-config is referenced from the corresponding task block in the main settings file. All paths inside sub-config files are defined relative to the `config_dir` specified in the common settings.

The *Simulation Config* and *Data Config* are reused across all three tasks and are therefore documented here to avoid repetition.

2.1 Simulation Config.json

The *Simulation Config* defines the simulation space and the wavefield propagation using the Beam Propagation Method (BPM).

Example: Simulation Config

```
{
  "Base_Grid": {
    "unit": "um",
    "wavelength": 0.640,
    "spatial_resolution": [0.1, 0.1, 0.1],
    "grid_shape": [250, 250, 250],
    "n_background": 1.334,
    "sim_mask": {
      "type": "sphere",
      "size": 110
    },
    "backpropagation_distance": 12.5
  }
}
```

Parameter Description

`unit`

Base unit used for Simulation Space and Propagation.

- "m", "mm", "um"

Note: All other unit will be converted to this unit.

`wavelength`

Wavelength of the propagating wavefield, expressed in `unit`.

`spatial_resolution`

Physical voxel size of the simulation grid, expressed in `unit`. Specified as $[d_x, d_y, d_z]$.

`grid_shape`

Number of voxels along each spatial dimension of the simulation grid.
Specified as $[n_x, n_y, n_z]$.

`n_background`

Base background refractive index of the simulation space.

`sim_mask`

Mask to limit the backwards mapping of the voxel object into the simulation space.

- {}: No mask applied.
- `type`:
 - `null`: No mask applied.
 - `sphere`: Spherical mask with radius `size`
 - `rect`: Rectangular mask with edges of length `size`

backpropagation_distance

Distance (in unit) over which the wavefield is propagated backwards after the forward BPM step, in order to reach the focal plane.

2.2 Data Config.json

The *Data Config* specifies the locations of all external data required by the different tasks. Paths are defined relative to the location of the executing script.

Not all data paths are required for every task.

Example: Data Config

```
{
  "Data": {
    "voxel_object": "../Data/Voxel Data/HEK Cells/DHM_Tomog.pt",
    "ground_truth": "../Data/GT Data/HEK Cells/DHM/DHM_Frames.csv",
    "mask": "../Data/Mask Data/HEK Cells/DHM_Mask.pt"
  }
}
```

Parameter Description

voxel_object

Path to the 3D refractive-index distribution.

Required for: {Forward Simulation, Pose Optimization, Reconstruction Optimization, Combined Optimization}

ground_truth

Path to a CSV file containing ordered ground-truth frame locations.

Required for: {Pose Optimization, Reconstruction Optimization, Combined Optimization}

mask

Path towards 3D Mask file, for masking out voxel object gradients

Optional for: {Reconstruction Optimization, Combined Optimization}

3 Forward Simulation

Usage

Given a set of key-frame poses, the estimated refractive-index (RI) distribution (*voxel object*) is placed into the simulation space using interpolated poses. A Beam Propagation Method (BPM) simulation is then performed to obtain the scattered wavefield.

A schematic overview of this process is shown in Fig. 3.

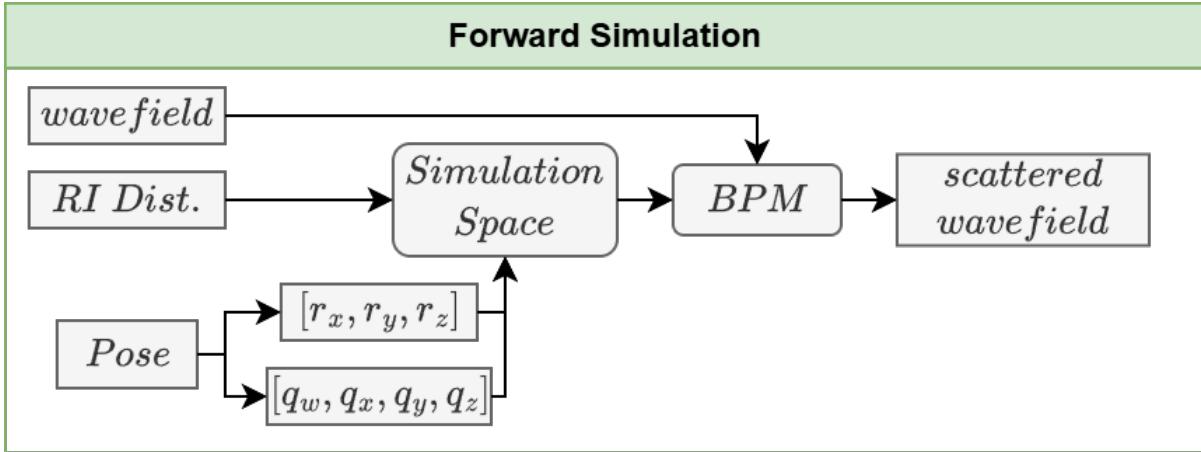


Figure 3: Schematic overview of the forward simulation process.

Script

- Forward_Simulation.py --settings "path/to/Settings.json"

Required Data

- *Settings* (.json)
 - Simulation Config (.json)
 - Data Config (.json)
 - Forward Config (.json)
- *Voxel object* (.pt): Estimated RI distribution

3.1 Configs

3.1.1 General Settings

Example: Settings.json

```
{
  "Forward_Simulation": {
    "simulation_config_file": "DHM/dhm_simulation.json",
    "data_config_file": "DHM/DHM_data.json",
    "forward_config_file": "Debug/Sim_Movements/debug_move.json",

    "output": {
      "output_dir": "../Outputs/Forward_Simulation/HEK_Cells/DHM/CleanUp",
      "options": ["phase", "amplitude", "slice", "sim_space", "file"]
    }
  }
}
```

Parameter Description

simulation_config_file

Path to the simulation configuration file defining the simulation space and wave propagation.

data_config_file

Path to the data configuration file containing the voxel object.

forward_config_file

Path to the configuration file defining object motion and post-processing operations.

output

see Sec. 3.2.

- **output_dir**: Output directory (created automatically)
- **options**: Output and visualization options
Options: {file, phase, amplitude, slice, sim_space}

3.1.2 Forward Config.json

The *Forward Config* defines object motion via poses specified at discrete key frames. Intermediate poses are obtained by interpolation.

If only a single key frame is provided, the pose is assumed to remain constant throughout the entire simulation. If the final key frame occurs before the last simulation timestep, the pose at that key frame is held constant until the end.

In addition, the Forward Config specifies post-processing operations applied independently to the wave-field, amplitude, and phase components. Details on available transforms are provided in Sec. 7.

Example: Forward Config

```
{
  "transforms": {
    "field": {
      "crop": {"x_min":125, "y_min":125, "x_max":-125, "y_max":-125},
      "upscale_complex": {"target_shape": [500,500]}
    },
    "amp": {},
    "phase": {
      "subtract_mean": {}
    }
  },
  "Movement": {
    "time_steps": 5,
    "unit": "um",
    "Positions": [
      { "pos": [0, 0, 0], "time": 0 },
      { "pos": [0, 0, 0], "time": 1 },
      { "pos": [2, -2, 0], "time": 2 }
    ],
    "Offset": [
      { "offset": [0, 0, 0], "time": 0 }
    ],
    "Rotation": [
      { "axis": [1, 0, 0], "theta": 0, "time": 0 },
      { "axis": [1, 0, 0], "theta": 5, "time": 1 },
      { "axis": [1, 1, 0], "theta": 45, "time": 4 }
    ]
  }
}
```

Parameter Description

Post-processing Transforms

transforms

See [7](#) for more information

- **field**: post-processing applied to the scattered wavefield
- **amp**: post-processing applied to the amplitude component.
- **phase**: post-processing applied to the phase component.

Movement**time_steps**

Total number of simulation timesteps.

unit

Unit for position, offset

Options: {"m", "mm", "um"}

Positions

Object center positions at specified key frames.

[{"pos" : [px, py, pz], "time" : ti}, ...]

Offset

Translation vector shifting the object center. Position and rotation are w.r.t. the object center.

[{"offset" : [ox, oy, oz], "time" : ti}, ...]

Rotation

Object rotation specified by axis/angle (deg) representation at key frames.

[{"axis" : [ax, ay, az], "theta" : theta_i, "time" : ti}, ...]

3.2 Outputs

Depending on the selected output **options**, various outputs are generated for each timestep. All outputs are stored in **output_dir/run_name**, which is created automatically if it does not exist. The output folder contains three subfolders: {Data, Images, Videos}.

Base Outputs

- **configs**: saves the simulation configs in separate json files.

Option: "file"

If "file" is in options, a .pt file is saved for each timestep containing:

wavefield

Complex scattered wavefield (only field post-processing)

amp

Amplitude image (includes amplitude post-processing)

phase

Phase image (includes phase post-processing)

pose_unit

Unit of pose parameters

position

Object position

offset

Object offset

axis

Rotation axis

angle

Rotation angle (degrees)

transforms

Applied post-processing transforms

sim_unit

Simulation-space unit

grid_shape

Simulation grid size

spatial_resolution

Simulation grid resolution

wavelength

Incident wavelength

Option: "amplitude"

If "amplitude" is in options, we save:

- Amplitude images for every timestep: `np.abs(field)`
- Amplitude video over all timesteps.

Option: "phase"

If "phase" is in options, we save:

- Phase images for every timestep: `unwrap_phase(np.angle(field))`
- Phase video over all timesteps.

Option: "slice"

If "slice" is in options, we save:

- Slice images of the simulation space at a specific index for every timestep.
- Slice video over all timesteps.

By default , the slice is taken at the center of the simulation space along the z-axis. These are non exposed parameters and can be modified in the `reporting.py` script if needed.

Option: "sim_space" or "render"

If "sim_space" or render is in options, we save:

- 3D rendering images of the simulation space for every timestep.
- 3D rendering video over all timesteps.

4 Pose Optimization

Usage

Given a recorded video sequence of scattered wavefields, the goal of pose optimization is to reconstruct the pose of the voxel object for every frame in the sequence.

This is achieved by optimizing the pose parameters of the *Forward Simulation* such that the simulated scattered wavefield best matches the recorded data.

We utilize a frame-by-frame approach, where we optimize each frame individually. Since video sequences are usually temporally smooth, we utilize the best pose of the previous frame as initialization for the next frame, to decrease the number of iterations until convergence.

A high-level overview of the frame-by-frame strategy is shown in Fig. 4, while the per-frame optimization workflow is illustrated in Fig. 5.

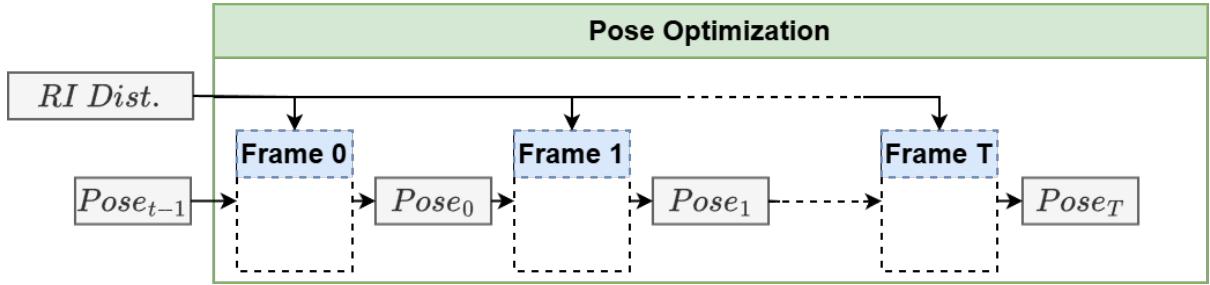


Figure 4: Frame-by-frame pose optimization overview.

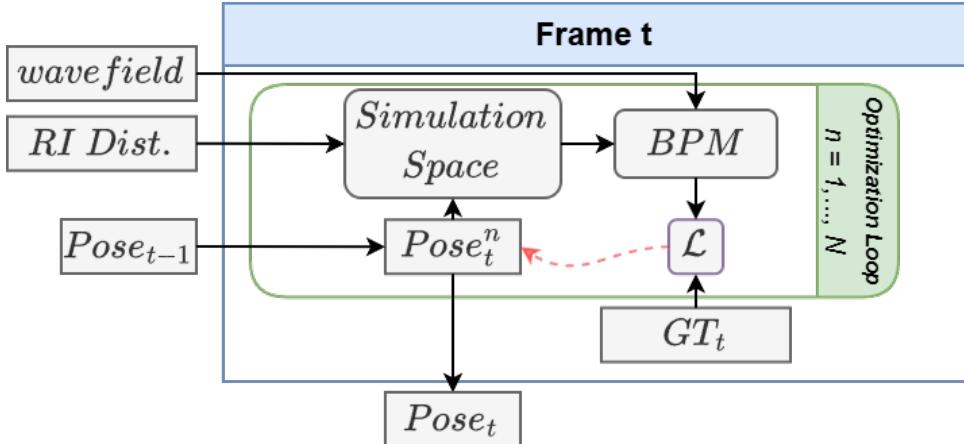


Figure 5: Per-frame pose optimization workflow.

Script

- Pose Optimization.py --settings "path/to/Settings.json"

Required Data

- **Settings file** (Settings.json)
 - Simulation Config
 - Data Config
 - PoseOpt Config
- **Voxel object** (.pt): Estimated RI distribution
- **Target data**: Recorded scattered wavefield sequence

4.1 Configs

4.1.1 General Settings

Example: Settings.json

```
{
  "PoseOpt": {
    "simulation_config_file": "DHM/dhm_simulation.json",
    "data_config_file": "DHM/DHM_data.json",
    "pose_opt_config_file": "DHM/DHM_pose_opt.json",

    "output": {
      "active": true,
      "output_dir": "../Outputs/PoseOpt/HEK Cells/DHM/Sample",
      "options": ["losses", "amps", "phases", "slices", "renders"]
    }
  }
}
```

Parameter Description

`simulation_config_file`

Path towards the config file defining the general simulation.

`data_config_file`

Path to the data configuration file containing voxel object and target data locations.

`pose_opt_config_file`

Path to the pose-optimization configuration file.

`output`

Output configuration (see Sec. 4.2).

- `active`: Enable or disable logging
- `output_dir`: Output directory (created automatically)
- `options`: Visualization and logging options
Options: {amps, phases, slices, renders}

4.1.2 PoseOpt Config.json

The *PoseOpt Config* specifies the optimization procedure used to estimate object poses over a sequence of frames. We handle the first frame within a sequence separately from all the other sequential frames. Therefore, we divide our PoseOpt Config into a general *Training*, a *Initial* frame, and *Sequential* frame section.

Example: PoseOpt Config

```
{
  "PoseOpt": {
    "gt_transforms": {
      "field": {},
      "amp": {
        "sqrt": {}
      },
      "phase": {
        "gaussian_blur": {"sigma": 7.5},
        "subtract_mean": {}
      }
    },
    "sim_transforms": {
      "field": {
        "upscale_complex": {"target_shape": [500, 500]}
      },
      "amp": {}
    }
  }
}
```

```

    "phase": {
        "subtract_mean": {}
    },
    "weights": [20, 75, 0.3, 300],
    "start_frame": 0,
    "end_frame": null,
    "frame_steps": 1,
    "unit": "um",
    "Position": [0.5, 0.25, 0],
    "Axis": [1, 0, 0],
    "Angle": 0,
    "Initial": {
        "epochs": 200,
        "optimizer": {
            "Position": {"lr": 0.01},
            "Quaternion": {"lr": 0.005}
        },
        "scheduler": {
            "milestones": [133, 166],
            "gamma": 0.1
        },
        "regularizers": {}
    },
    "Sequential": {
        "epochs": 20,
        "optimizer": {
            "Position": {"lr": 0.001},
            "Quaternion": {"lr": 0.0005}
        },
        "scheduler": {
            "milestones": [30],
            "gamma": 0.5
        },
        "regularizers": {
            "Position": {"lambda": 1, "w": [0.01, 0.01, 1], "v": [0.01, 0.01, 1]},
            "Quaternion": {"lambda": 50, "w": [1e-5, 1e-5, 1e-3, 1e-3], "v": [1e-5, 1e-5, 1e-3, 1e-3]}
        }
    }
}
}

```

Parameter Description

Training

gt_transforms

Domain-adaptation transforms applied to the recorded data (see Sec. 7).

sim_transforms

Domain-adaptation transforms applied to the simulated wavefield (see Sec. 7).

weights

Component-wise weights for the loss function.

start_frame, end_frame, frame_steps

Frame indices defining the optimization range and sampling.

unit

Unit used for positional parameters.

Position, Axis, Angle

Initial pose parameters for the first frame.

Initial Frame**epochs**

Number of optimization epochs for the first frame.

optimizer

Adam Optimizer for the first frame

- **Position, Quaternion**
 - lr: Learning rate for the respective parameter

scheduler

MultiStepLR scheduler for the first frame

- **milestones**: Epochs at which to decay the learning rate
- **gamma**: Decay factor

regularizers

By default, no regularization is applied for the initial frame.

Sequential Frames**epochs**

Number of optimization epochs per subsequent frame.

optimizer

Adam Optimizer configuration for sequential frames.

- **Position, Quaternion**
 - lr: Learning rate for the respective parameter

scheduler

MultiStepLR scheduler for sequential frames.

- **milestones**: Epochs at which to decay the learning rate
- **gamma**: Decay factor

regularizers

Kalman Filter regularization for sequential frames

- **Position**
 - lambda: Regularization strength
 - w: Process noise covariance (float or 3-element list)
 - v: Measurement noise covariance (float or 3-element list)
- **Quaternion**
 - lambda: Regularization strength
 - w: Process noise covariance (float or 4-element list)
 - v: Measurement noise covariance (float or 4-element list)

4.2 Outputs

Depending on the selected output options, pose optimization produces per-frame outputs as well as a summary of the whole sequence. All outputs are written to `output_dir/run_name`, which is created automatically if it does not exist.

The directory contains three subfolders: `{Summary, Configs, Frames}`.

Base Outputs

Configs

configs: Simulation configuration files (JSON).

- Settings.json
- Simulation.json
- Data.json
- PoseOpt.json

Frames

Images: Best settings visualizations for the current frame.

- Best Amplitude.png
- Best Phase.png
- Best Slice.png
- Best Render.png

Summary

Images: Position, quaternion, axis, and angle plots over all frames.

- best_positions.png
- best_quaternions.png
- best_axes.png
- best_angles.png

Videos: Best amplitude, phase, slice, and render videos over all frames.

- Best Amplitude.avi
- Best Phase.avi
- Best Slice.avi
- Best Render.avi

Option: "losses"

Frames

Losses: Total/Data/Reg Loss and individual loss components per epoch in a csv file.

- Losses.csv

Plot: Loss plot (total, data, and regularization) for current frame.

- Frame Loss.png

Summary

Plot: Plots Total/Data/Reg Loss and individual loss components of best settings over all frames.

- Losses.png

Best Epochs: Plot of best iteration index for each frame.

- best_epochs.png

Option: "amps"

Frames

amplitude Periodically saves and amplitude for current settings in amp subfolder.

- amp/Amplitude_Epoch_xxx.png

Option: "phases"**Frames**

phase Periodically saves and phase for current settings in phases subfolder.

- `phase/Phase_Epoch_xxx.png`

Option: "poses"**Frames**

Poses: Saves position and rotation values for all iterations of current frame in csv files.

- `Positions.csv`
- `Rotations.csv`

5 Reconstruction Optimization

Usage

Given a recorded video sequence of scattered wavefields and a list of poses, reconstruction optimization updates the estimated refractive-index (RI) distribution of the voxel object such that the simulated wavefields generated by the *Forward Simulation* better match the recorded data.

The voxel object is updated only after processing the full dataset; that is, one optimization step corresponds to a complete pass through all recorded frames.

A schematic overview of the reconstruction process is shown in Fig. 6.

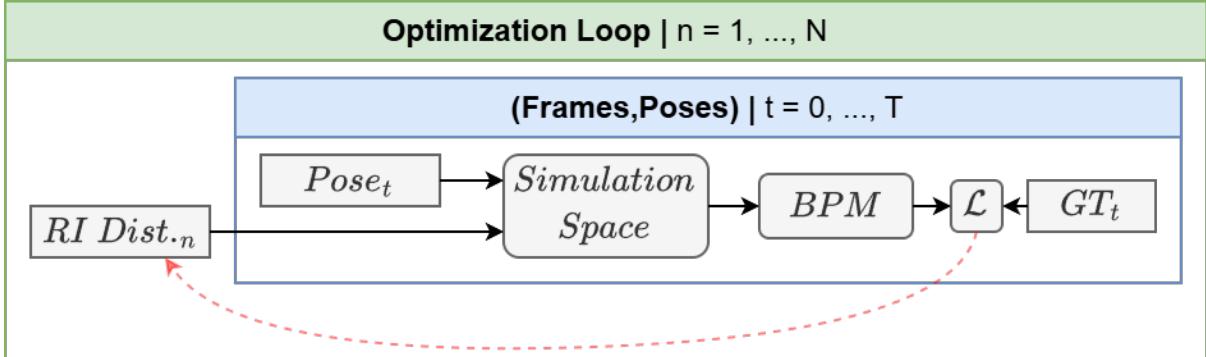


Figure 6: Reconstruction optimization workflow.

Script

- Reconstruction Optimization.py --settings "path/to/Settings.json"

Required Data

- **Settings file** (Settings.json)
 - Simulation Config
 - Data Config
 - ReconOpt Config
- **Voxel object** (.pt): Initial RI distribution
- **Poses** (.json): i.e. output of Pose Optimization
- **Target data**: Recorded scattered wavefield sequence

Optional

- **Gradient mask** (.pt): Mask to restrict voxel object updates

5.1 Configs

5.1.1 General Settings

Example: Settings.json

```
{
  "ReconOpt": {
    "simulation_config_file": "DHM/dhm_simulation.json",
    "data_config_file": "DHM/DHM_data.json",
    "recon_opt_config_file": "DHM/dhm_recon_opt.json",
    "recon_poses_file": "Debug/Optimized Poses/dhm_poses_sample.json",

    "output": {
      "active": true,
    }
  }
}
```

```
        "output_dir": "../Outputs/ReconOpt/HEK Cells/DHM/Sample",
        "options": ["slices", "renders", "amps", "phases"]
    }
}
```

Parameter Description

simulation_config_file

Path to the simulation configuration file.

data_config_file

Path to the data configuration file containing voxel object and target data locations.

`recon_opt_config_file`

Path to the reconstruction optimization configuration file.

recon_pose_file

Path to a JSON file containing optimized poses for all frames.

output

Output configuration (see Sec. 5.2).

- `active`: Enable or disable logging
 - `output_dir`: Output directory (created automatically)
 - `options`: Visualization options
 - Options:** {amps, phases, slices, renders}

5.1.2 ReconOpt Config.json

Example: ReconOpt Config

```

{
  "ReconOpt": {
    "gt_transforms": {
      "field": {},
      "amp": {
        "sqrt": {}
      },
      "phase": {
        "gaussian_blur": {"sigma": 7.5},
        "subtract_mean": {}
      }
    },
    "sim_transforms": {
      "field": {
        "upscale_complex": {"target_shape": [500, 500]}
      },
      "amp": {},
      "phase": {
        "subtract_mean": {}
      }
    },
    "epochs": 20,
    "weights": [20, 75, 0.3, 300],
    "optimizer": {
      "Voxel Object": {"lr": 1e-4}
    },
    "scheduler": {
      "milestones": [15, 18],
      "gamma": 0.1
    }
  }
}

```

```

    "regularizers" : {
        "TV Reg": {"lambda": 500}
    }
}
}
}

```

Parameter Description

gt_transforms

Domain-adaptation transforms applied to recorded data (see Sec. 7).

sim_transforms

Domain-adaptation transforms applied to simulated wavefields (see Sec. 7).

epochs

Number of training iterations. Each iteration corresponds to one full pass through the dataset.

weights

Component-wise weights for the loss function.

optimizer

Adam Optimizer configuration used to update the voxel object

- Voxel Object
 - lr: Learning rate for the respective parameter

scheduler

MultiStepLR scheduler configuration.

- milestones: Epochs at which to decay the learning rate
- gamma: Decay factor

regularizers

Total Variation regularization for the voxel object

- TV Reg:
 - lambda: Regularization strength

5.2 Outputs

Depending on the selected output options, reconstruction optimization produces outputs after each full dataset pass. All outputs are written to `output_dir/run_name`, which is created automatically if it does not exist.

The directory contains three subfolders: `{Summary, Configs, Epochs}`.

Base Outputs

Configs

Copies of all configuration files used during execution.

- `Settings.json`
- `Simulation.json`
- `Data.json`
- `ReconOpt.json`
- `ReconPoses.json`

Summary

Wavefield Images: Amplitude/Phase of optimized voxel object with reconstruction poses

- `Amps/Amplitude_xxx.png`

- Phases/Phase_xxx.png

Voxel Object Images: Slice/Render of optimized voxel object at best epoch

- Slice_xxx.png
- Render_xxx.png

Wavefield Videos: Amplitude/Phase comparison of optimized voxel object with reconstruction poses

- Amplitude.avi
- Phase.avi

Extended Summary Plots: Combination of Reconstruction Poses and Amplitude/Phase Comparison

- Plots/Summary_xxx.png
- Summary.avi

Option: "slices"

Epochs

Slices: Slice of voxel object for current epoch.

- Slice_xxx.png

Summary

Voxel Object Progress Video: Video of optimization progress of voxel object slice over all epochs.

- Voxel Object Slice.avi

Option: "renders"

Epochs

Renders: 3D render of voxel object for current epoch.

- Render_xxx.png

Summary

Voxel Object Progress Video: Video of optimization progress of 3D render of voxel object over all epochs.

- Voxel Object Render.avi

Option: "amps"

Epochs

Image: saves amplitude image of last processed pose for current epoch.

- Amplitude_xxx.png

Option: "phases"

Epochs

Image: saves phase image of last processed pose for current epoch.

- Phase_xxx.png

Option: "terse"

Extended Plots of Combination of Reconstruction Poses and Amplitude/Phase Comparison (only speeds up final summary).

6 Combined Optimization

Usage

Combined optimization jointly optimizes both the object poses and the refractive-index (RI) distribution of the voxel object given a recorded video sequence of scattered wavefields.

This is achieved by alternating between pose optimization (see Sec. 4) and reconstruction optimization (see Sec. 5) for a set number of cycles.

A high-level overview of the combined optimization procedure is shown in Fig. 7.

At the moment the number of combined optimization cycles is fixed in the code itself (pose optimization followed by reconstruction optimization). Furthermore, Combined Optimization expects a *gaussian_blur* domain adaptation transform for the phase component in the *gt_transforms* of both the PoseOpt and ReconOpt configs. The sigma value of the blur is used to progressively decrease the amount of domain adaptation over the combined optimization cycles.

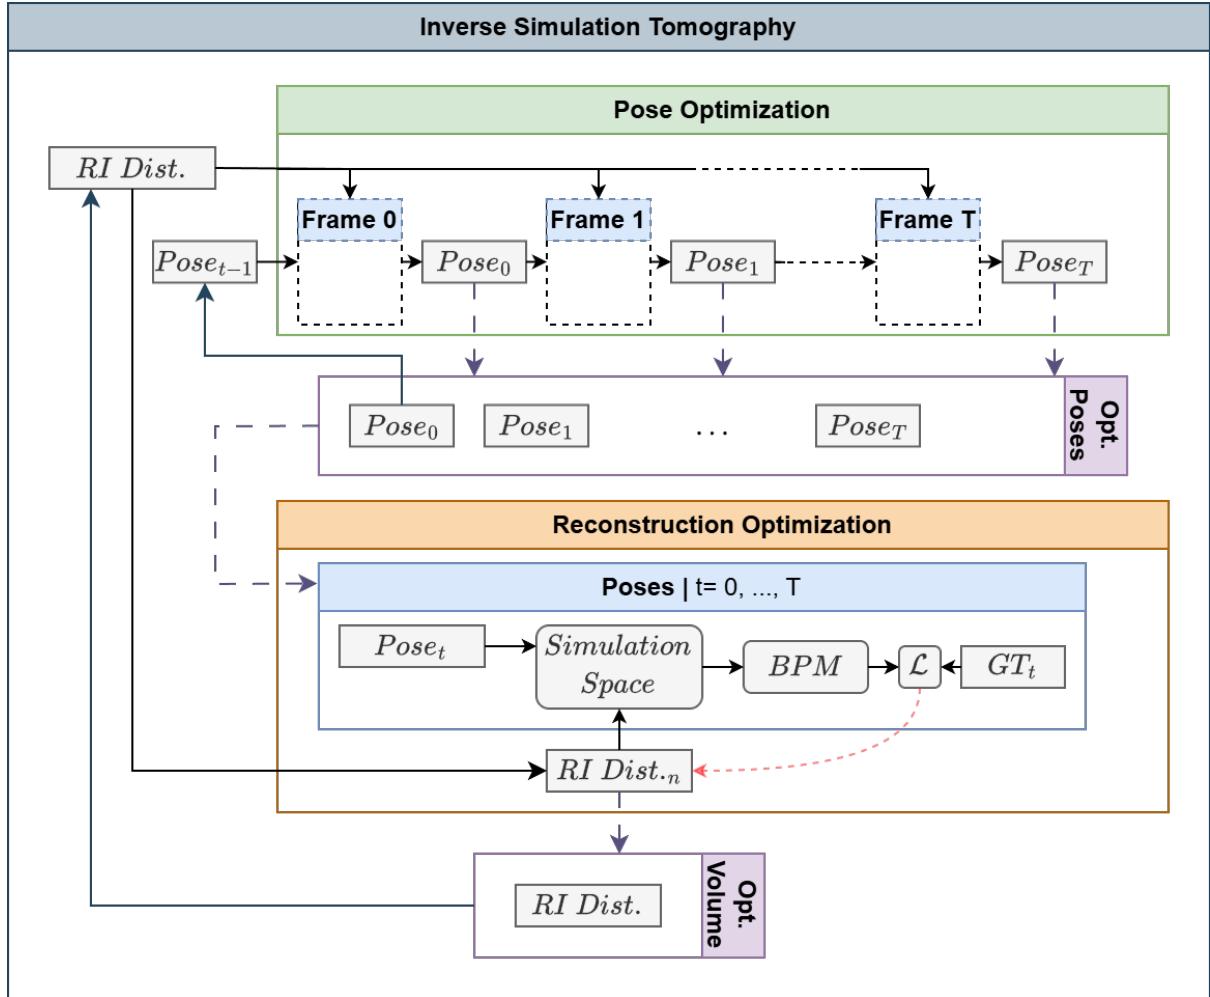


Figure 7: Combined optimization workflow.

Script

- `Combined Optimization.py --settings "path/to/Settings.json"`

Required Data

- **Settings file (Settings.json)**
 - Simulation Config

- Data Config
- PoseOpt Config
- ReconOpt Config
- **Voxel object (.pt)**: Initial RI distribution

Optional

- **Gradient mask (.pt)**: Mask to restrict voxel object updates

6.1 Configs

6.1.1 General Settings

Example: `Settings.json`

```
{
  "CombOpt": {
    "simulation_config_file": "DHM/dhm_simulation.json",
    "data_config_file": "DHM/DHM_data.json",
    "pose_opt_config_file": "DHM/dhm_pose_opt.json",
    "recon_opt_config_file": "DHM/dhm_recon_opt.json"

    "output": {
      "output_dir": "../Outputs/CombOpt/HEK Cells/DHM/12_12_25 Long + Always Update",
      "options": ["slices", "renders", "amps", "phases"]
    }
  }
}
```

Parameter Description

`simulation_config_file`

Path to the simulation configuration file.

`data_config_file`

Path to the data configuration file containing voxel object and target data locations.

`pose_opt_config_file`

Path to the pose-optimization configuration file.

`recon_opt_config_file`

Path to the reconstruction-optimization configuration file.

`output`

Output configuration (see Sec. 6.2).

- `output_dir`: Output directory (created automatically)
- `options`: Visualization options
Options: {amps, phases, slices, renders}

6.2 Outputs

For every combined optimization cycle, outputs from both the PoseOpt and ReconOpt procedures are saved in separate subfolders within `output_dir/run_name`.

The directory contains a subfolder for each cycle named: `iter xxx`, where `xxx` denotes the cycle index (starting from 0).

Within each cycle subfolder, outputs from both optimization procedures are stored in their respective subfolders.

7 Domain Adaptation / Post-Processing Transforms

Transforms serve two main purposes:

1. Post-processing simulation outputs.
2. Domain adaptation during training, increasing similarity between simulated and recorded data.

Conceptually, transforms are divided into three components: `field`, `amp`, and `phase`. Each component defines an ordered list of functions applied sequentially. Since the amplitude and phase are derived from the complex wavefield after applying `field` transforms, any such transforms will directly affect them.

A schematic pipeline for domain adaptation is shown in Fig. 8.

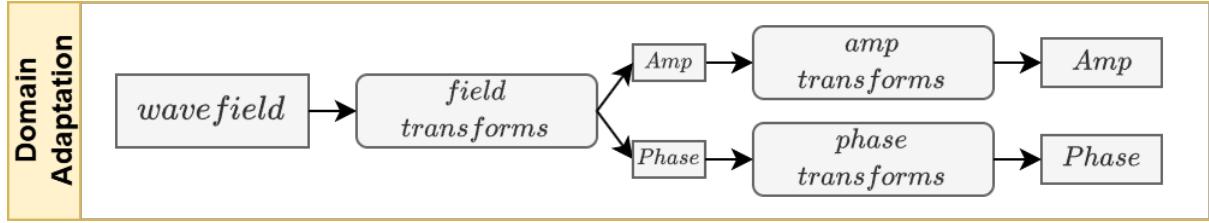


Figure 8: Domain adaptation pipeline showing transform application.

Example

```
{
  "transforms": {
    "field": {
      "crop": {"x_min":125, "y_min":125, "x_max":-125, "y_max":-125},
      "upscale_complex": {"target_shape": [500,500]}
    },
    "amp": {
      "sqrt": {}
    },
    "phase": {
      "subtract_mean": {}
    }
  }
}
```

Available Transform Functions

Each transform is defined by a `name` and a dictionary of parameters.

`crop`

Crops a rectangle defined by pixel coordinates.

`x_min`, `y_min` Starting pixel index.

`x_max`, `y_max` Ending pixel index.

`sqrt`

Computes the square root of the image values.

- No additional parameters.

`upscale_complex`

Bilinear upscale complex valued 2D tensor to a target size.

`target_shape` Tuple specifying the new image size.

`subtract_mean`

Subtracts the mean of the image.

- No additional parameters.

subtract_min

Shifts all pixel values to be positive by subtracting the minimum.

- No additional parameters.

subtract_background

Subtracts the mean of a small corner patch near the origin.

- No additional parameters.

gaussian.blur

Applies Gaussian blur.

sigma Standard deviation for Gaussian kernel.

gaussian.blur.scheduled

Applies Gaussian blur, decreasing over a number of steps.

max_sigma Starting sigma value.

min_sigma Ending sigma value.

num_steps Number of steps to decrease sigma over.

snr_noise

Adds noise based on the SNR.

snr_db Noise level in decibels.

noise

Adds Gaussian noise.

sigma Standard deviation of the noise.

8 Custom Data

Similar to the Configs files this project expects the data to be organized in a specific structure. This section describes how to create custom data for use with the simulation and optimization scripts.

8.1 Ground Truth Data

These are the recorded scattered wavefields that the simulations will try to match. We expect every frame to be stored as a separate .pt file.

To load the data our dataloader expects a CSV file containing the paths to each frame. The order of the paths in the CSV file defines the frame order.

Each .pt file should contain the following entries:

field_data_file

Path towards the source data

Currently: Unused.

frame_index

Frame index.

amp Amplitude values of the recorded wavefield.

opd Optical path difference values of the recorded wavefield.

opd_units

Units of the optical path difference values.

px_size

Pixel size of the recorded wavefield.

Currently: Unused

position

Estimated object position.

Currently: Unused

offset

Estimated object offset.

Currently: Unused

rotation_axis

Estimated Rotation axis.

Currently: Unused

angle

Estimated Rotation angle (degrees).

Currently: Unused

8.2 Voxel Object

This is the Initial Voxel Object used for the simulations. It should be stored as a .pt file containing:

data

3D tensor containing the refractive index distribution.

spatial_resolution

Voxel size

unit

Unit of spatial resolution

is_lefthanded

Bool indicating if the coordinate system is left-handed. (Simulation uses Left-handed system)

8.3 Gradient Masks

Optionally we can use a gradient mask to mask out certain regions of the voxel object during reconstruction optimization. The gradient mask should be stored as a .pt file containing:

`mask`

3D tensor containing values between 0 and 1.